# Predicting Water Pump Failures in Tanzania

## A Machine Learning Project by

Austine Denis Otieno

## Phase3

# Overview

- Rural Tanzanians rely heavily on water pumps.
- Many pumps fail without early warnings, disrupting access to clean water.
- My goal: Use data and machine learning to predict pump failures and help decision-makers prioritize maintenance.

# Business and Data Understanding

- Dataset from Taarifa and DrivenData: 59,000plus water points.
- Problem: Predict the pump's status:  Functional, Needs Repair, or Non-Functional.
- Business value: Enables proactive maintenance, reducing outages and improving public health.

# The Data

- **Features include:**
  - **GPS coordinates, region, installer**
  - **Water quality, pump type, extraction method**
  - **Construction year, quantity of water, etc.**
- **Challenges:**
  - **Many categorical variables**
  - **Imbalanced class distribution (most pumps functional)**

# Data Preparation

- Cleaned missing data and standardized inputs
- Combined rare categories into 'Other' to simplify patterns
- Applied SMOTE to address class imbalance
- One-hot encoding used for categorical features

# Modeling

- Tried several models: Logistic Regression, Decision Trees, Random Forest
- Random Forest performed best:
  - Handles categorical and numerical data well
  - Robust to overfitting

# Logistic regression

Logistic Regression Accuracy: 0.5941077441077441

| | precision | recall | f1-score | support |
|---|---|---|---|---|
| 0 | 0.59 | 0.89 | 0.71 | 6452 |
| 1 | 0.00 | 0.00 | 0.00 | 863 |
| 2 | 0.60 | 0.29 | 0.39 | 4565 |
| accuracy | | | 0.59 | 11880 |
| macro avg | 0.40 | 0.39 | 0.37 | 11880 |
| weighted avg | 0.55 | 0.59 | 0.54 | 11880 |

Class 0 has a lot more samples than class 1 → the model learns to predict class 0 often.

Random Forest Accuracy: 0.8068181818181818

| | precision | recall | f1-score | support |
|---|---|---|---|---|
| 0 | 0.81 | 0.89 | 0.85 | 6452 |
| 1 | 0.52 | 0.32 | 0.40 | 863 |
| 2 | 0.84 | 0.78 | 0.81 | 4565 |
| accuracy | | | 0.81 | 11880 |
| macro avg | 0.72 | 0.66 | 0.69 | 11880 |
| weighted avg | 0.80 | 0.81 | 0.80 | 11880 |

Overall accuracy is up to 80.8%—a big jump from 59%.

Accuracy: 0.8084175084175084

| | precision | recall | f1-score | support |
|---|---|---|---|---|
| 0 | 0.79 | 0.92 | 0.85 | 6452 |
| 1 | 0.62 | 0.24 | 0.35 | 863 |
| 2 | 0.85 | 0.76 | 0.81 | 4565 |
| accuracy | | | 0.81 | 11880 |
| macro avg | 0.76 | 0.64 | 0.67 | 11880 |
| weighted avg | 0.80 | 0.81 | 0.80 | 11880 |

The model strongly predicts class 0 and class 2, which together make up the majority of the dataset.

# Evaluation

- Used Accuracy, Precision, Recall, F1-Score
- Final Accuracy: 81%
- High performance on Functional and Non-Functional classes
- Moderate performance on Needs Repair (due to few examples)

# Interpretation of Results

- **Class 0 (Functional):** Most predictions accurate
- **Class 1 (Needs Repair):** Often confused with other classes
- **Class 2 (Non-functional):** High recall and precision
- Indicates value for prioritizing emergency interventions

# Recommendations:

- Use model to highlight pumps at high risk of failure
- Improve field data quality and consistency
- Update model regularly with new data
- Train local staff to interpret and act on predictions

# Next Steps

- Integrate model with mobile or web-based reporting tools
- Collaborate with government agencies for deployment
- Explore use of satellite or weather data to enhance predictions

# THANK YOU